

Matrix Transpose 2D


Parallel Algorithm Design

Outline


1. Transpose Algorithm
 - a. Tiling
 - b. Cache Oblivious Algorithm
2. Memory Benchmark
 - a. Media Results
 - b. Rome Results
3. Computational Benchmark
 - a. Media Results
 - b. Rome Results

Transpose Algorithm

$a[j * N + 0]$







$b[0 * N + j]$



- switch rows and columns
→ memory bound
→ stride access
- no use of locality in naive version
→ solution: tiling

Tiling



- use of spatial locality
- tiles must fit in cache
- additional control flow instructions



Cache-Oblivious Transpose

TBB: `blocked_range2d` + `simple_partitioner`

- termination condition defined by grain size



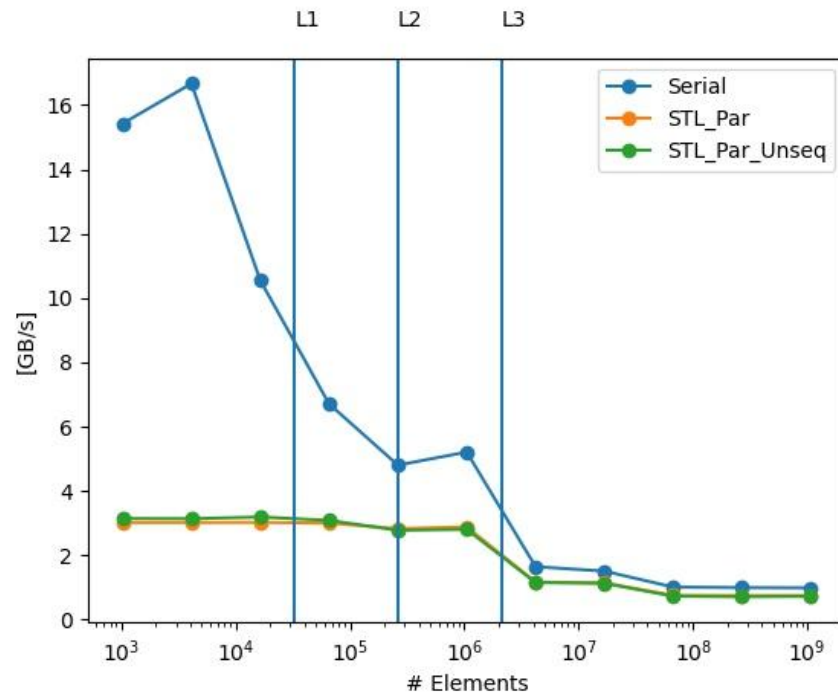
→ traversal pattern

Used Hardware:

| | mp-media (Intel Xeon E3-1585 v5) | ziti-rome (AMD Epyc 7662) |
|---|--|--|
| base frequency | 3900 MHz | 2154.3 MHz |
| NUMA nodes | 1 | 4 |
| cores | 4 physical 8 logical | 64 physical 128 logical |
| cache sizes | 4x 32KB L1 4x 256KB L2 1x 8MB L3 | 64x 32KB L1 64x 512KB L2 16x 16MB L3 |
| memory bandwidth (read & write concurrently) | 18.2 GB/s | 98.9 GB/s |

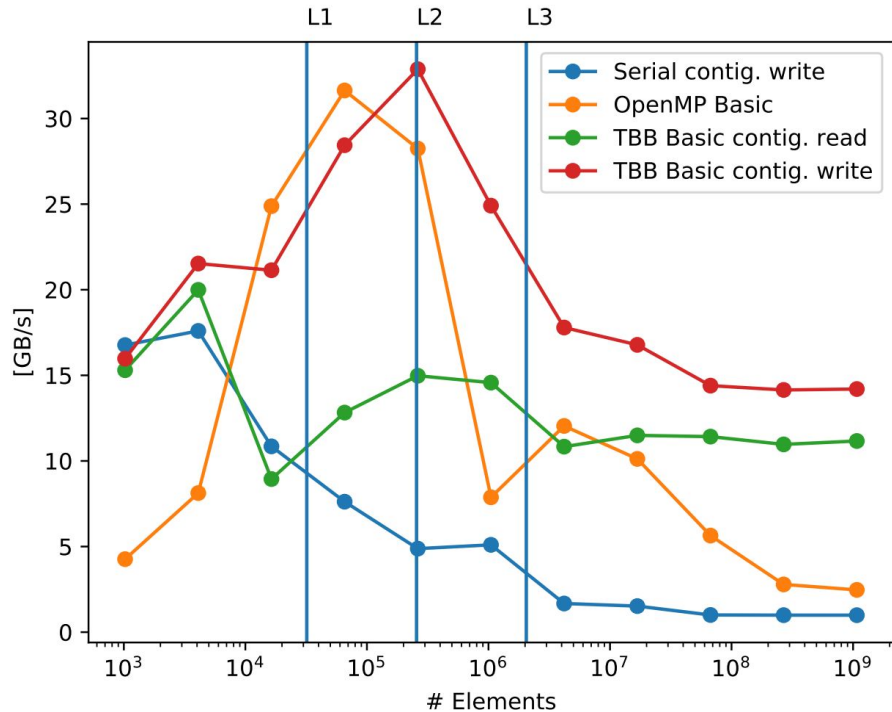
Memory Benchmark STL

- Transposition with `std::for_each()` and parallel execution policy
- leads to very bad results
- `std::for_each()` hard to use for matrix transposition



Memory Benchmark Media


- TBB:
 - simple partitioner
 - grain size: 64
- OMP:
 - dynamic scheduling
- basic TBB version performs better
 - use of locality



Memory Benchmark Media Tiling

- TBB:
 - simple partitioner
 - grain size: 64
- OMP:
 - dynamic scheduling
 - block size: 64
 - speedup of 5.7 to OMP basic
- no effect by using OMP SIMD

75% peak mem. performance




Memory Benchmark Media AVX Intrinsic

- TBB:
 - simple partitioner
 - avx with grain size: 128
- OMP:
 - dynamic scheduling
 - avx with block size: 128

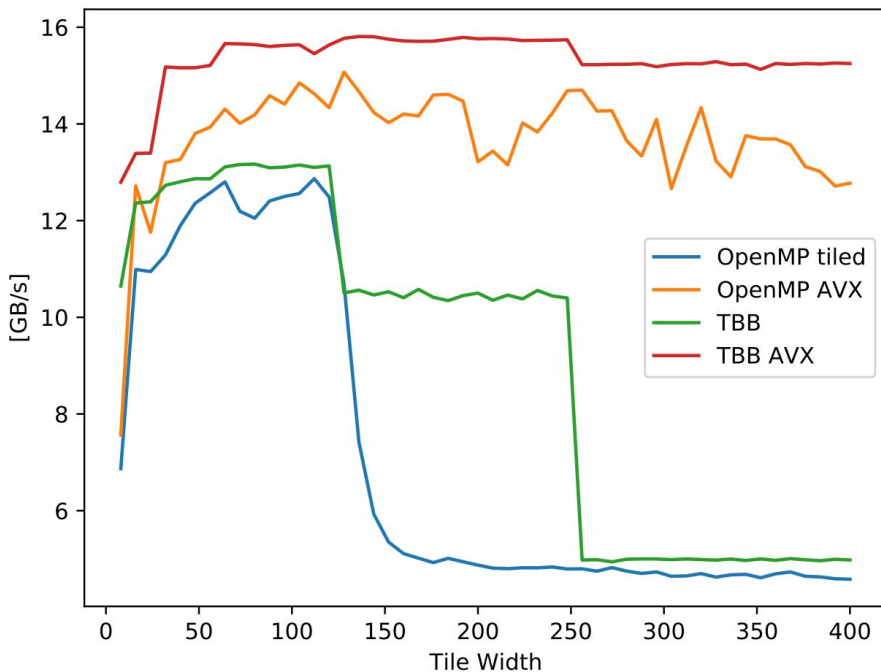
87% peak mem. performance

(non avx versions tile size 64)



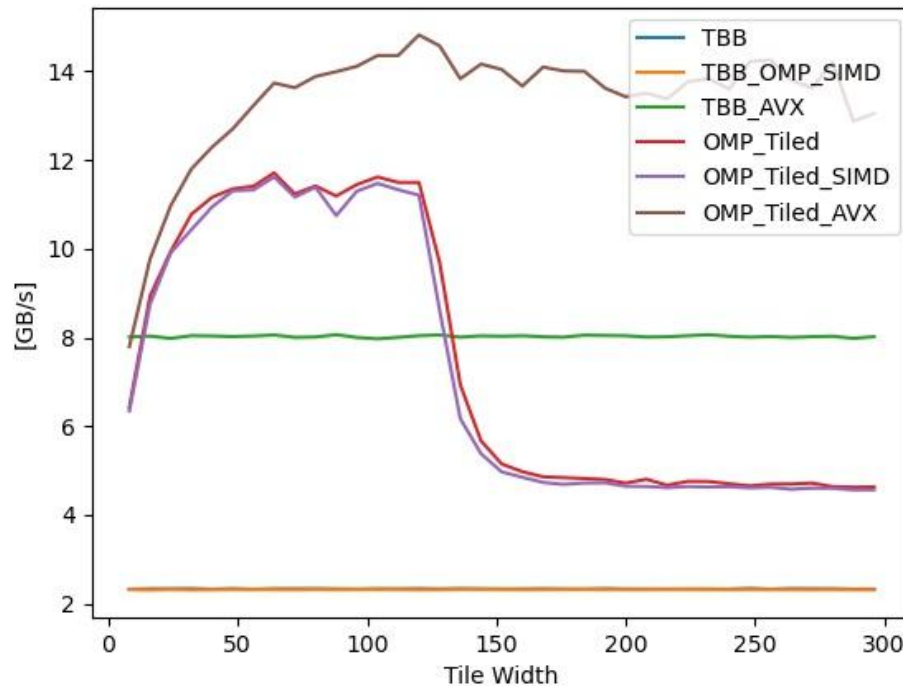
Memory Benchmark Media Optimal Tile Size

- OMP tiled version extremely bad with wrong tile size
- TBB performs worse when limiting traversal pattern to big tiles
- AVX versions quite stable
 - because of second layer of tiling
 - 8x8 blocks with 256bit registers



Memory Benchmark Media Static Scheduling

- static partitioner with TBB breaks tiling
→ no control over tile size
- static scheduling for OMP slightly worse performance



Memory Benchmark Rome

- with best media config

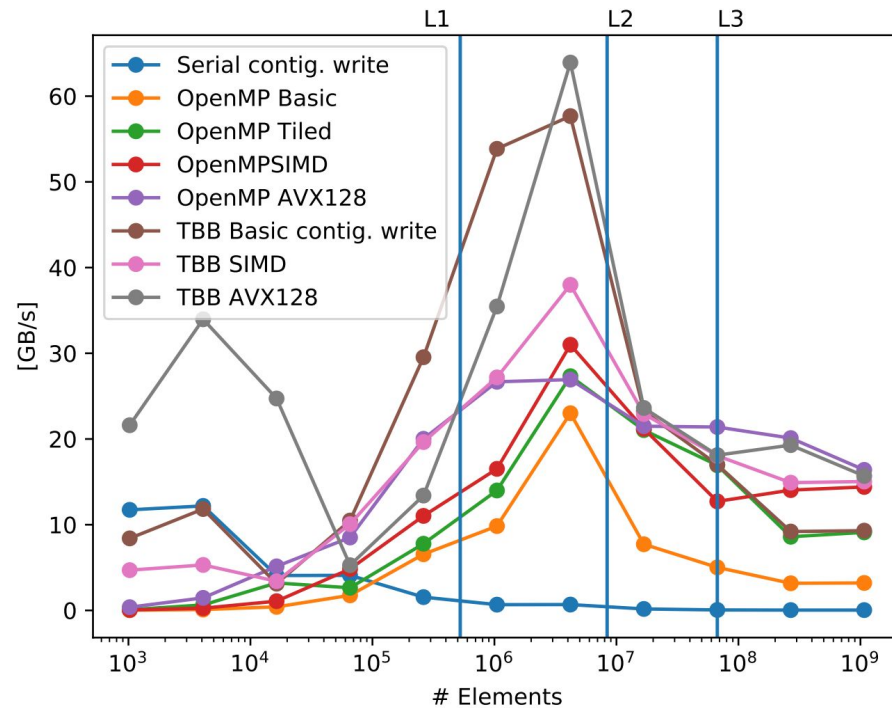
TBB:

- simple partitioner
- avx with grain size: 128

OMP:

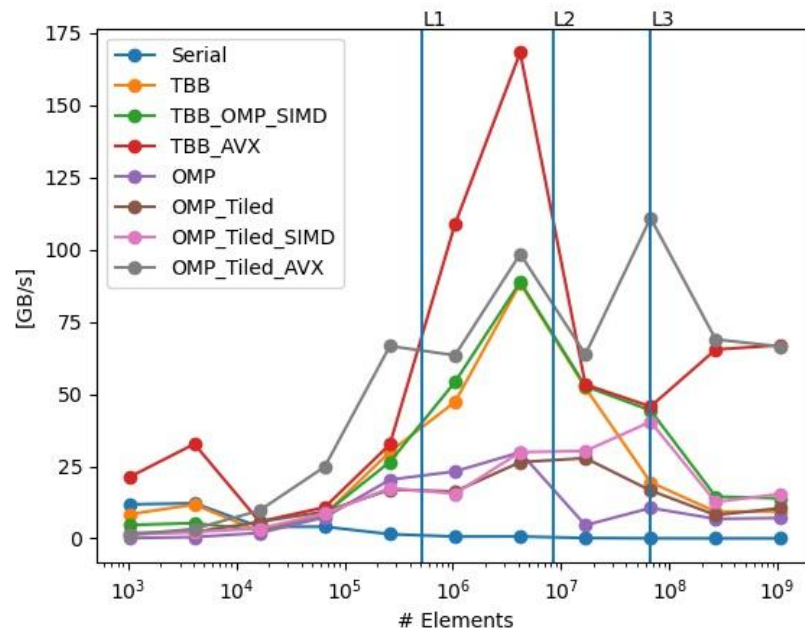
- dynamic scheduling
- avx with block size: 128

- bad performance
→ NUMA effects



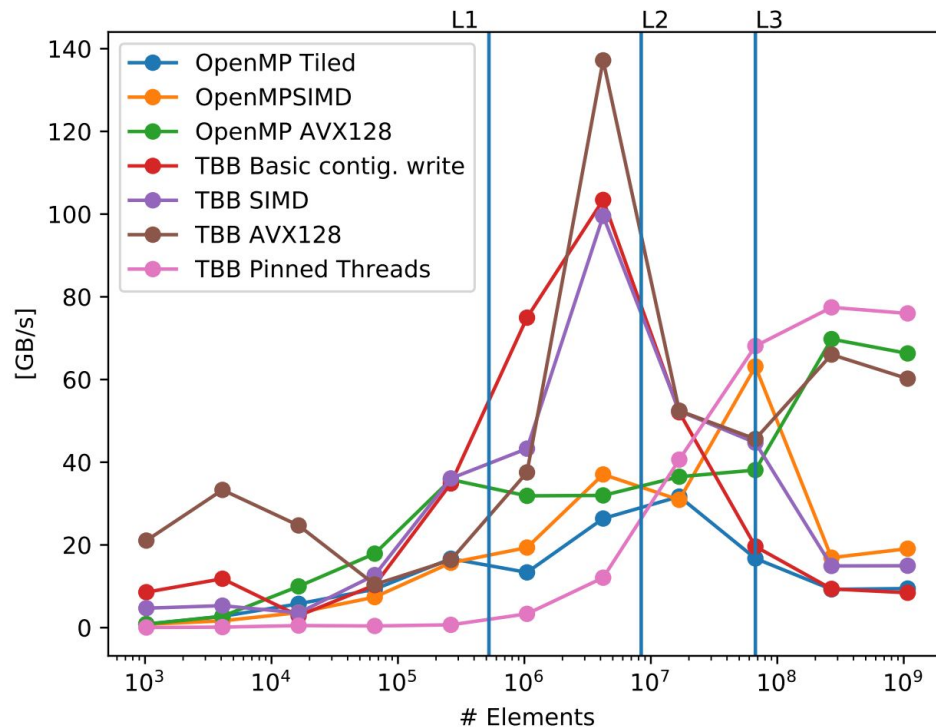
Memory Benchmark Rome

- NUMA aware data placement
(static scheduling for OMP)
- TBB still simple partitioner
- non intrinsic versions still bad




Memory Benchmark Rome

- additionally use of OMP Places
→ threads locked to cores
→ slight improvement
- for TBB version with hwloc
→ matrix divided in 4 parts
→ computed separately in NUMA regions
- wrong tile size on new machine?



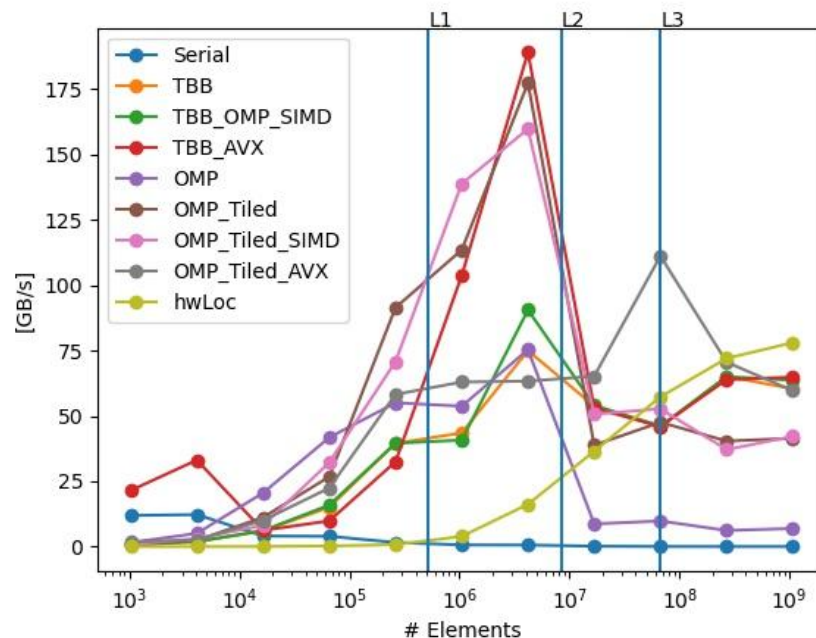
Memory Benchmark Rome Optimal Tile Size

- matrix size fixed to $2^{15} \times 2^{15}$
- small tiles perform very good
→ 8 x 8 looks promising



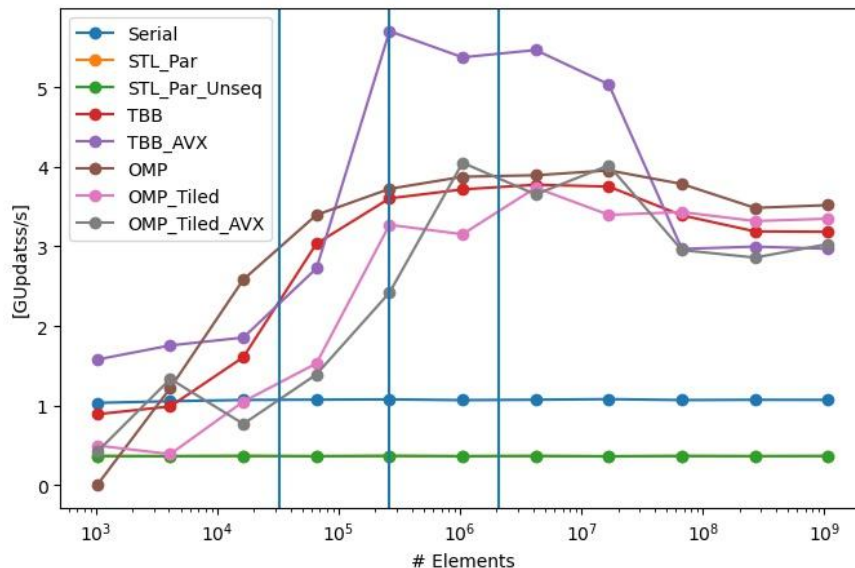
Memory Benchmark Rome Optimal Tile Size

- significant performance increase with 8x8 tiles
- AVX versions still use 64x64 tiles
but: internal 8x8 tiling
- 80% peak performance with hwloc (third layer of tiling)



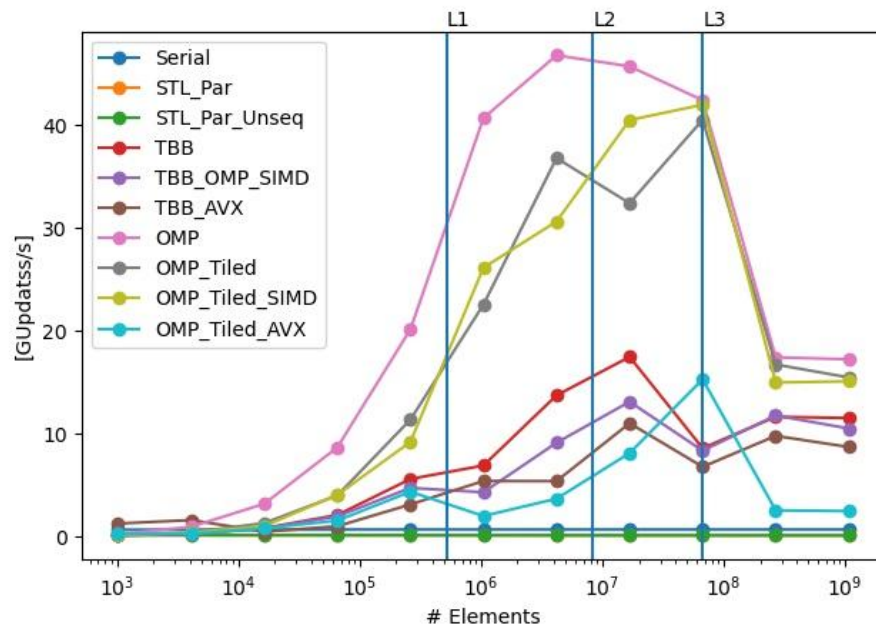
Computational Media

- same parameters from memory benchmark
- basic OMP version best
 - no stride in input matrix
 - contiguous write in output matrix
- overhead in AVX version
→ fill phase of local arrays



Computational Rome

- same parameters from memory benchmarks (8x8 blocks)
- OMP Basic still best performance
→ maximum write performance



Thanks for your attention